

Introduction to Choreographies

Fabrizio Montesi

`fmontesi@imada.sdu.dk`

Department of Mathematics and Computer Science

University of Southern Denmark

Last update: November 7, 2018

This document contains lecture notes on choreographies for the course on DM861 Concurrency Theory (2018) at the University of Southern Denmark. It is organised as a working book. Different parts get updated during the course.

The chapters in advance of the lectures are given as a preview on what is to come, but **remember to download the updated version of this book every week to stay up to date!**

Contents

<i>Preface: Alice, Bob, and Choreographies</i>	5
<i>This book</i>	9
1 Inference systems	11
1.1 Example: Flight connections	11
1.2 Derivations	13
1.3 Proof non-existence	15
1.4 Rule derivability and admissibility	18
1.4.1 Derivable rules	19
1.4.2 Rule admissibility	20
2 Simple choreographies	23
2.1 Syntax	24
2.2 Semantics	24
3 From choreographic programs to process programs	27
3.1 Simple processes	28
3.1.1 Syntax	29
3.1.2 Semantics: discussion	29
3.1.3 Semantics: definition	31
3.2 Endpoint Projection (EPP)	34
3.2.1 From choreographies to processes	34
3.2.2 Towards a correct EPP	36
3.2.3 Simple Concurrent Choreographies	38
4 Statefulness	41
4.1 Local Computation	41
4.1.1 Stateful Choreographies	41
4.1.2 Stateful Processes	45
4.1.3 EndPoint Projection	45

CONTENTS

A Solution to selected exercises	47
List of Notations	51
Index	52

Preface:

Alice, Bob, and Choreographies

We live in the era of concurrency and distribution. Concurrency gives us the ability to perform multiple tasks at a time. Distribution allows us to do so at a distance. It is the era of the web, mobile computers, cloud computing, and microservices. The number of computer programs that communicate with other programs over a network is exploding. By 2025, the Internet alone might be expected to connect from 25 to 100 billion devices [OECD, 2016].

Modern computer networks and their applications are key drivers of our technological advancement. They give us better citizen services, a more efficient industry (Industry 4.0), new ways to connect socially, and even better health with smart medical devices.

Modern computer networks and their applications are complex. Services are becoming increasingly dependent on other services. For example, a web store might depend on an external service provided by a bank to carry out customer payments. The web store, the customer's web browser, and the bank service are thus integrated: they communicate with each other to reach the goal of transferring the right amount of money to the right recipient, such that the customer can get the product she wants from the store. In distributed systems, the heart of integration is the notion of *protocol*: a document that prescribes the communications that the involved parties should perform in order to reach a goal.

It is important that protocols are clear and precise, because each party needs to know what it is supposed to do such that integration is successful and the whole system works. At the same time, it is important that protocols are as concise as possible: the bigger a protocol, the higher the chance that we made some mistake in writing it. Computer scientists and mathematicians might get a familiar feeling when presented with the necessity of achieving clarity, precision, and conciseness in writing. A computer scientist could point out that we need a good *language* to write protocols. A mathematician might say that we need a good *notation*.

Needham and Schroeder [1978] introduced an interesting notation for writing protocols. A communication of a message M from the participant A to the participant B is written

$$A \rightarrow B : M.$$


To define a protocol where A sends a message M to B , and then B passes the same message to C , we can just compose communications in a sequence:

$$\begin{aligned} A &\rightarrow B : M \\ B &\rightarrow C : M. \end{aligned}$$

This notation is called “Alice and Bob notation”, due to a presentational style found in security research whereby the participants A and B represent the fictional characters Alice and Bob, respectively, who use the protocol to perform some task. There might be more participants, like C in our example—typically a shorthand for Carol, or Charlie. The first mention of Alice and Bob appeared in the seminal paper by Rivest et al. [1978] on their public-key cryptosystem:

“For our scenarios we suppose that A and B (also known as Alice and Bob) are two users of a public-key cryptosystem”.

Over the years, researchers and developers created many more protocol notations. Some of these notations are graphical rather than textual, like Message Sequence Charts [International Telecommunication Union, 1996]. The message sequence chart of our protocol for Alice, Bob, and Charlie looks as follows.



For our particular example, the graphical representation of our protocol (as a message sequence chart) and our previous textual representation (in Alice and Bob notation) are equivalent, in the sense that they contain the

same information. This is not the case in general: not all notations are equally expressive, as we will see in the rest of this book.¹

In the beginning of the 2000s, researchers and practitioners took the idea of protocol notations even further, and developed the idea of *choreography*. A choreography is essentially still a protocol, but the emphasis is on detail. Choreographies might include details like the following.

- The kind of data being transmitted, or even the actual functions used to compute the data to be transmitted.
- Explicit cause-effect relations (also known as causal dependencies) between the data being transmitted and the choices made by participants in a protocol.
- The state of participants, e.g., their memory states.

Choreographies are typically written in languages designed to be readable *mechanically*, which also make them amenable to be used by a computer. In this book, we will start with a very simple choreography language and then progressively extend it with more sophisticated features, like parallelism and recursion. We will see that it is possible to define mathematically a *semantics* for choreographies, which gives us an interpretation of what running a protocol means. We will also see that it is possible to translate choreographies to a theoretical model of executable programs, which gives us an interpretation of how choreographies can be correctly implemented in the real world.

Although choreographies still represent a young and active area of research, they have already started emerging in many places. In 2005, the World Wide Web Consortium (W3C)—the main international standards organisation for the web—drafted the Web Services Choreography Description Language (WS-CDL), for defining interactions among web services. In 2011, the global technology standards consortium Object Management Group (OMG) introduced choreographies in their notation for business processes [BPMN]. The recent paradigm of microservices [Dragoni et al., 2017] advocates for the use of choreographies to achieve better scalability. All this momentum is motivating a lot of research on both the theory of choreographies and its application to programming [Ancona et al., 2016, Hüttel et al., 2016].

It is the time of Alice and Bob. It is the time of choreographies.

¹Message sequence charts in particular support many features, including timeouts and alternative behaviours.

This book

This book is an introduction to the theory of choreographies. It explains what choreographies are and how we can model them mathematically. Its primary audiences are computer science students and researchers. However, the book should be approachable by anybody interested in studying choreographies, e.g., for theoretical purposes or the development of choreography-based tools. The aim of this book is to be pedagogical. It is not an aim to be comprehensive. References to alternative notations and techniques to model choreographies are given where appropriate. It is assumed that the reader is familiar with the notion of concurrency and how distributed systems are programmed.

Prerequisites To read this book, you should be familiar with:

- discrete mathematics and the induction proof method;
- context-free grammars (only basic knowledge is required);
- basic data structures, like trees and graphs;
- concurrent and distributed systems.

These prerequisites are attainable in most computer science B.Sc. degrees, or with three years of relevant experience.

To study choreographies, we are going to define choreography languages and then write choreographies as terms of these languages. The syntax of languages is going to be defined in terms of context-free grammars. To give meaning to choreographies, we are going to use extensively Plotkin's structural approach to operational semantics.

The rules defining the semantics of choreographies are going to be rules of inference, borrowing from deductive systems. Knowing formal systems based on rules of inference is not a requirement to read this book: chapter 1 provides an introduction to the essential knowledge on these systems that we need for the rest of the book. The reader familiar with inference systems or

structural operational semantics can safely skip the first chapter and jump straight to chapter 2.

An important aspect of choreographies is determining how they can be executed correctly in concurrent and distributed systems, in terms of independent programs for *processes*. To model process programs, we will borrow techniques from the area of process calculi. We will introduce the necessary notions on process calculi as we go along, so knowing this area is not a requirement for reading this book. The reader familiar with process calculi will recognise that we borrow ideas from Milner’s seminal calculus of communicating systems [Milner, 1980].

Chapter 1

Inference systems

Before we venture into the study of choreographies, we need to become familiar with the formalism that we are going to use throughout this book: inference systems. Inference systems are widely used in formal logic and the specification of programming languages (our case).¹

An inference system is a set of *inference rules* (also called rules of inference). An inference rule has the form

$$\frac{\text{Premise 1} \quad \text{Premise 2} \quad \cdots \quad \text{Premise } n}{\text{Conclusion}}$$


and reads “If the premises Premise 1, Premise 2, \dots , and Premise n hold, then Conclusion holds”. It is perfectly valid for an inference rule to have no premises: we call such rules *axioms*, because their conclusions always hold. An inference rule has always exactly one conclusion.

1.1 Example: Flight connections

This example is inspired by the usage of rules of inference to model graphs by Pfenning [2012].

Consider the following undirected graph of direct flights between cities.

¹For further reading on these systems, see the lecture notes by Martin-Löf [1996].



Let A , B , and C range over cities. We denote that two cities A and B are connected by a direct flight with the *proposition* $\text{conn}(A, B)$. Then, we can represent our graph as the set of axioms below.

$$\frac{}{\text{conn}(\text{Odense}, \text{Munich})} \quad \frac{}{\text{conn}(\text{Munich}, \text{Rome})} \quad \frac{}{\text{conn}(\text{Paris}, \text{Munich})}$$

$$\frac{}{\text{conn}(\text{Paris}, \text{Lisbon})}$$

Notice that our graph is undirected: in an ideal world, all direct flights are available also in the opposite direction. This is not faithfully represented by our axioms: if $\text{conn}(A, B)$, it should also be the case that $\text{conn}(B, A)$. One option to solve this discrepancy is to double the number of our axioms, to include their symmetric versions—e.g., for $\text{conn}(\text{Munich}, \text{Odense})$). A more elegant option is to include a rule of inference for symmetry, as follows.

$$\frac{\text{conn}(A, B)}{\text{conn}(B, A)} \text{SYM}$$

The label SYM is the name of our new inference rule; it is just a decoration to remember what the rule does (SYM is a shorthand for symmetry). Rule SYM tells us that if we have a connection from any A to any B (premise),

1.2. Derivations

$$\begin{array}{c}
 \overline{\text{conn}(\text{Odense, Munich})} \quad \overline{\text{conn}(\text{Munich, Rome})} \quad \overline{\text{conn}(\text{Paris, Munich})} \\
 \\
 \overline{\text{conn}(\text{Paris, Lisbon})} \\
 \\
 \frac{\text{conn}(A, B)}{\text{conn}(B, A)} \text{SYM} \quad \frac{\text{conn}(A, B)}{\text{path}(A, B)} \text{DIR} \quad \frac{\text{path}(A, B) \quad \text{path}(B, C)}{\text{path}(A, C)} \text{TRANS}
 \end{array}$$

Figure 1.1: An inference system for flights.

then we have a connection from B to A (conclusion). In this rule, A and B are *schematic variables*: they can stand for any of our cities.

Now that we are satisfied with how our rule system captures the graph, we can use it to find flight paths from a city to another. For any two cities A and B , the proposition $\text{path}(A, B)$ denotes that there is a path from A to B . Finding paths is relatively easy. First, we observe that direct connections give us a path. (DIR stands for direct.)

$$\frac{\text{conn}(A, B)}{\text{path}(A, B)} \text{DIR}$$

Second, we formulate a rule for multi-step paths. If there is a path from A to B , and a path from B to C , then we have a path from A to C . In other words, paths are transitive. (TRANS in the following rule stands for transitivity.)

$$\frac{\text{path}(A, B) \quad \text{path}(B, C)}{\text{path}(A, C)} \text{TRANS}$$

The whole system is displayed in fig. 1.1.

1.2 Derivations

The key feature of an inference system is the capability of performing *derivations*. Suppose that we wanted to answer the following question.

Is there a flight path from Odense to Rome?

Answering this question in our inference system for flight connections corresponds to showing that $\text{path}(\text{Odense, Rome})$ holds. To do this, we build a *derivation tree* (also simply called *derivation*, or *proof tree*). The problem tackled in the following is also known as proof search.

We start our derivation from what we want to conclude with (our conclusion is what we want to prove).

$\text{path}(\text{Odense}, \text{Rome})$

Observe that we have only two inference rules that can conclude something of this form: **DIR** and **TRANS**. **DIR** cannot be applied: we would need to prove $\text{conn}(\text{Odense}, \text{Rome})$, which is impossible. So our only choice is rule **TRANS**, by instantiating A as Odense and C as Rome . How should we instantiate B in rule **TRANS**, though?

$$\frac{\text{path}(\text{Odense}, ??B??) \quad \text{path}(??B??, \text{Rome})}{\text{path}(\text{Odense}, \text{Rome})} \text{TRANS}$$

We need to guess a correct instantiation for B and hope that it will allow us to continue our derivation. Looking at the definition of our graph, it is easy to see that the right choice is to pick Munich .

$$\frac{\text{path}(\text{Odense}, \text{Munich}) \quad \text{path}(\text{Munich}, \text{Rome})}{\text{path}(\text{Odense}, \text{Rome})} \text{TRANS}$$

Our derivation is not complete yet, because we are left with the tasks of proving (deriving) $\text{path}(\text{Odense}, \text{Munich})$ and $\text{path}(\text{Munich}, \text{Rome})$. Thanks to the notation of inference rules, we can just “dig deeper” with further rule applications. Since we know that Odense and Munich are connected, we can try using rule **DIR**.

$$\frac{\frac{\text{conn}(\text{Odense}, \text{Munich})}{\text{path}(\text{Odense}, \text{Munich})} \text{DIR} \quad \text{path}(\text{Munich}, \text{Rome})}{\text{path}(\text{Odense}, \text{Rome})} \text{TRANS}$$

We now have to prove $\text{conn}(\text{Odense}, \text{Munich})$. We have that as the conclusion of one of our axioms, so we can conclude that branch of our derivation by applying the related axiom.

$$\frac{\frac{\text{conn}(\text{Odense}, \text{Munich})}{\text{path}(\text{Odense}, \text{Munich})} \text{DIR} \quad \text{path}(\text{Munich}, \text{Rome})}{\text{path}(\text{Odense}, \text{Rome})} \text{TRANS}$$

We can finish our derivation for the right premise $\text{path}(\text{Munich}, \text{Rome})$ likewise.

1.3. Proof non-existence

$$\frac{\frac{\text{conn}(\text{Odense}, \text{Munich})}{\text{path}(\text{Odense}, \text{Munich})} \text{DIR} \quad \frac{\frac{\text{conn}(\text{Munich}, \text{Rome})}{\text{path}(\text{Munich}, \text{Rome})} \text{DIR}}{\text{path}(\text{Odense}, \text{Rome})} \text{TRANS}$$

Our derivation is done! We can tell by the fact that there are no more premises left to prove.

If you look carefully, you can see that our derivation is a tree. The conclusion $\text{path}(\text{Odense}, \text{Rome})$ is the root of the tree, and the leaves are empty nodes (the premises of our axioms). Rule applications connect the nodes of the tree. In fact, even our previous partial derivations are all trees. This is a very useful property that we will use throughout the book.

We impose that derivations are finite: we are interested in proofs that can be checked mechanically in finite time.

1.3 Proof non-existence

*Proof, or proof of no proof. There is no try.
- Yoda, to an exhausted Luke*

Showing that a proposition holds requires showing a proof, i.e., a derivation, in the inference systems of interest. Once the proof is shown, then we just need to check that all rules have been applied correctly. If we are convinced that this is the case, then we are convinced that the proof is correct and that the proposition indeed holds.

What about showing that a proposition *cannot* be derived? This can be far trickier, because it requires us to reason about all the possible proofs that could, potentially, conclude with our proposition and showing that none of those proofs can actually be built.

Consider a simple example: showing that $\text{conn}(\text{Lisbon}, \text{Rome})$ is not derivable. Intuitively, there is no direct connection between Lisbon and Rome in our graph. But how can we show this formally?

First, we observe that the only rules that can have a conclusion of the form $\text{conn}(A, B)$ for any A and B are our axioms and rule SYM. None of the axioms can be applied for $A = \text{Lisbon}$ and $B = \text{Rome}$, as in our case. We are left with rule SYM: any search of a derivation of $\text{conn}(\text{Lisbon}, \text{Rome})$ must begin as follows.

$$\frac{\text{conn}(\text{Rome}, \text{Lisbon})}{\text{conn}(\text{Lisbon}, \text{Rome})} \text{SYM}$$

By similar reasoning (there is no rule to apply for `conn`(Rome, Lisbon) but `SYM`), we obtain that the proof *must* continue with another application of `SYM`.

$$\frac{\frac{\text{conn}(\text{Lisbon}, \text{Rome})}{\text{conn}(\text{Rome}, \text{Lisbon})} \text{SYM}}{\text{conn}(\text{Lisbon}, \text{Rome})} \text{SYM}$$

We got back to where we started: we have to prove `conn`(Lisbon, Rome). Since we have only one way to prove it, and we have just shown that it leads to the exact same premise, this points out that our proof search will go on indefinitely and that we will never reach a finite derivation. So, `conn`(Lisbon, Rome) cannot be proven.

Let us be more formal, to convince ourselves that `conn`(Lisbon, Rome) cannot be derived more decisively. Since every derivation is a finite tree, we can measure the height of a derivation as a natural number (it is the height of the tree). Thus, among all the proofs of `conn`(Lisbon, Rome), there is at least one of minimal height, in the sense that there is no other proof of `conn`(Lisbon, Rome) with lower height. We attempt at finding this minimal proof. The reasoning goes as before, and we soon end up seeing that a minimal proof necessarily starts as follows.

$$\frac{\frac{\text{conn}(\text{Lisbon}, \text{Rome})}{\text{conn}(\text{Rome}, \text{Lisbon})} \text{SYM}}{\text{conn}(\text{Lisbon}, \text{Rome})} \text{SYM}$$

So we have to find some proof of our premise `conn`(Lisbon, Rome) on top. But if such a proof exists, it would be a proof of `conn`(Lisbon, Rome) that is *smaller* than the proof that we are building (because it would not have the first two applications of `SYM` of our proof). Thus our minimal proof must be bigger than another proof, and we reach a contradiction.

Another way of proving that `conn`(Lisbon, Rome) is by looking at how proofs are constructed from the top, rather than the bottom. For our system, we can prove the following result. We use \mathcal{P} to range over proofs (derivations).

Proposition 1. *For any natural number n , there exists no proof of `conn`(Lisbon, Rome) of height n .*

Proof. We prove the stronger result that there exists no proof of `conn`(Lisbon, Rome) and there exists no proof of `conn`(Rome, Lisbon), either.

We proceed by induction on n .

1.3. Proof non-existence

$$\begin{array}{c}
\overline{\text{conn}(\text{Odense}, \text{Munich})} \quad \overline{\text{conn}(\text{Munich}, \text{Rome})} \quad \overline{\text{conn}(\text{Paris}, \text{Munich})} \\
\\
\overline{\text{conn}(\text{Paris}, \text{Lisbon})} \\
\\
\frac{\text{conn}(A, B)}{\text{conn}(B, A)} \text{SYM} \\
\\
\frac{\text{conn}(A, B)}{\text{path}(A, B, 1)} \text{DIRW} \quad \frac{\text{path}(A, B, n) \quad \text{path}(B, C, m)}{\text{path}(A, C, n + m)} \text{TRANSW}
\end{array}$$

Figure 1.2: Weighted rules for flight paths.

Base case: $n = 1$. All proofs with height 1 must necessarily be an application of an axiom, and there is no axiom that can prove either $\text{conn}(\text{Lisbon}, \text{Rome})$ or $\text{conn}(\text{Rome}, \text{Lisbon})$.

Inductive case: $n = m + 1$ for some natural number m . By induction hypothesis, we know that there is no proof \mathcal{P} such that the height of \mathcal{P} is m and that the conclusion of \mathcal{P} is $\text{conn}(\text{Lisbon}, \text{Rome})$ or $\text{conn}(\text{Rome}, \text{Lisbon})$. Thus, all proofs \mathcal{Q} of height m conclude with either: i) $\text{conn}(A, B)$ for some A and B such that the set $\{A, B\}$ is different from $\{\text{Lisbon}, \text{Rome}\}$; or ii) $\text{path}(A, B)$ for some A and B . For i), we observe that there is no rule that allows us to derive $\text{conn}(\text{Lisbon}, \text{Rome})$ from a premise that is not $\text{conn}(\text{Rome}, \text{Lisbon})$, and likewise for the symmetric case where the premise is $\text{conn}(\text{Rome}, \text{Lisbon})$. For ii), we observe that there is no rule that, given a conn proposition as one of its premises, allows us to conclude $\text{conn}(\text{Lisbon}, \text{Rome})$ or $\text{conn}(\text{Rome}, \text{Lisbon})$.

□

It follows from proposition 1 that there is no proof of $\text{conn}(\text{Lisbon}, \text{Rome})$. Assume that there were such a proof. Since it would have to be finite, it would have a height n . Thus we reach a contradiction, because proposition 1 states that for all n there is no proof of $\text{conn}(\text{Lisbon}, \text{Rome})$.

Exercise 1. Consider the system in fig. 1.2, which replaces rules DIR and TRANS with alternative rules that measure the length of a path (paths are “weighted”, with each connection having weight 1).

$$\begin{array}{c}
 \overline{\text{conn}(\text{Odense, Munich})} \quad \overline{\text{conn}(\text{Munich, Rome})} \quad \overline{\text{conn}(\text{Paris, Lisbon})} \\
 \\
 \frac{\text{conn}(A, B)}{\text{conn}(B, A)} \text{SYM} \\
 \\
 \frac{\text{conn}(A, B)}{\text{path}(A, B, 1)} \text{DIRW} \quad \frac{\text{path}(A, B, n) \quad \text{path}(B, C, m)}{\text{path}(A, C, n + m)} \text{TRANSW}
 \end{array}$$

Figure 1.3: A limited and weighted flight system.

$$\begin{array}{c}
 \overline{\text{conn}(\text{Odense, Munich})} \quad \overline{\text{conn}(\text{Munich, Rome})} \quad \overline{\text{conn}(\text{Paris, Munich})} \\
 \\
 \overline{\text{conn}(\text{Paris, Lisbon})} \\
 \\
 \frac{\text{conn}(A, B)}{\text{conn}(B, A)} \text{SYM} \quad \frac{\text{conn}(A, B)}{\text{path}(A, B)} \text{DIR} \quad \frac{\text{conn}(A, B) \quad \text{path}(B, C)}{\text{path}(A, C)} \text{STEP}
 \end{array}$$

Figure 1.4: An alternative way of constructing paths.

Prove that, for any A and B , if $\text{path}(A, B)$ is derivable in the system in fig. 1.1, then there exists n such that $\text{path}(A, B, n)$ is derivable in the system in fig. 1.2.

Suggestion: proceed by structural induction on the proof of $\text{path}(A, B)$.

Exercise 2 (!). *Consider the system in fig. 1.3, which removes the direct flight from Paris to Munich.*

Prove that it is not possible to derive $\text{path}(\text{Lisbon, Munich}, n)$ for any n .

1.4 Rule derivability and admissibility

Consider the system in fig. 1.4, which replaces rule TRANS from fig. 1.1 with rule STEP. The difference is that rule STEP requires a direct connection from the source A to some city B , and then a path from B to the destination C .

From an algorithmic perspective, adopting rule TRANS or rule STEP might lead to slightly different search strategies. Searching for a path from A to C using rule STEP roughly corresponds to: look up in our database of

1.4. Rule derivability and admissibility

direct connections (given by the axioms and their reflexive closure, thanks to rule SYM) from A to some B , and then recursively try to find a path from B to C ; if we fail, we have to try with another B , if possible. By contrast, searching for a path from A to C using rule TRANS corresponds to recursively trying to find a path from A to some B , and then again recursively trying to find a path from B to C ; again, if we fail, we have to try with another B , if possible. Of course, these strategies are only for paths not covered already by rule DIR, which covers the case in which A and C are connected directly.

Since the two systems are different, a key question is whether they are *equally powerful*, in the sense that every derivable proposition in one of the two is derivable also in the other.

There are only two kinds of propositions that we can derive in our two systems: $\text{conn}(A, B)$ and $\text{path}(A, B)$. It is easy to see that, for any A and B , $\text{conn}(A, B)$ is derivable in the system with rule TRANS if and only if it is derivable in the system with rule STEP, simply because the two systems share exactly the same rules for deriving conn propositions. For propositions of the form $\text{path}(A, B)$, the situation is more complicated. We tackle the two directions separately (from the system with rule STEP to the system with rule TRANS, and vice versa).

1.4.1 Derivable rules

We start by showing that the system with rule TRANS (fig. 1.1) can derive all paths that can be derived in the system with rule STEP (fig. 1.4).

To do this, we prove that adding rule STEP to the system with rule TRANS would not add any new derivable propositions. Recall that rule STEP looks as follows.

$$\frac{\text{conn}(A, B) \quad \text{path}(B, C)}{\text{path}(A, C)} \text{ STEP}$$

The key observation here is that we can build a derivation that, from the premises $\text{conn}(A, B)$ and $\text{path}(B, C)$, concludes $\text{path}(A, C)$ by using the rules in fig. 1.1. Here it is.

$$\frac{\frac{\text{conn}(A, B)}{\text{path}(A, B)} \text{ DIR} \quad \text{path}(B, C)}{\text{path}(A, C)} \text{ TRANS}$$

The derivation above is proof that rule STEP is *derivable* in the system in fig. 1.1. In general, we say that a rule is derivable whenever its conclusion can be derived from its premises by using rules that are already in the system. In other words, if we can build a derivation from the premises of the rule to its conclusion, then the rule is derivable.

1.4.2 Rule admissibility

Let us look at the other direction: proving that if adding rule TRANS to the system with rule STEP would not add any new derivable propositions.

Recall that rule TRANS is defined as follows.

$$\frac{\text{path}(A, B) \quad \text{path}(B, C)}{\text{path}(A, C)} \text{ TRANS}$$

As a first attempt, we could try to adopt the same strategy that we followed in section 1.4.1: deriving the conclusion $\text{path}(A, C)$ from the premises $\text{path}(A, B)$ and $\text{path}(B, C)$ using the rules in fig. 1.4.

Unfortunately, we reach a dead end pretty quickly when trying to show that rule TRANS is derivable in the system with rule STEP. The only way to build a path with multiple connections is by using rule STEP, which requires a **conn** as premise. But our only available premises are **path** propositions, and we have no rule that allows us to conclude **conn** from a **path**.

We resort to a different proof technique and show that rule TRANS is *admissible* in the system with rule STEP (fig. 1.4). An admissible rule is one that does not add any new derivable propositions. All derivable rules are also admissible, but admissible rules are not necessarily derivable (just like our case here for rule TRANS). It is sometimes convenient to mark admissible rules explicitly when defining them. We will use dashed horizontal lines to denote admissible inference rules.

Theorem 1. *The rule*

$$\frac{\text{path}(A, B) \quad \text{path}(B, C)}{\text{path}(A, C)} \text{ TRANS}$$

is admissible in the system in fig. 1.4.

Proof. Let \mathcal{P} and \mathcal{Q} be the derivations of $\text{path}(A, B)$ and $\text{path}(B, C)$, respectively. We proceed by induction on the size of the derivation

$$\frac{\frac{\mathcal{P}}{\text{path}(A, B)} \quad \frac{\mathcal{Q}}{\text{path}(B, C)}}{\text{path}(A, C)} \text{ TRANS}$$

and prove that we can build a derivation using the rules in fig. 1.4 *that is not bigger* (this is going to be important for the last case of the proof). We have three cases, depending on the last applied rule in \mathcal{P} (one for each rule that can conclude a **path** proposition).

1.4. Rule derivability and admissibility

Case Dir \mathcal{P} ends with an application of rule **DIR**. Thus, for some \mathcal{P}' that does not contain any applications of rule **TRANS**, we are in the following situation.

$$\frac{\frac{\mathcal{P}'}{\text{conn}(A, B)} \quad \frac{\mathcal{Q}}{\text{path}(B, C)}}{\text{path}(A, C)} \text{TRANS}$$

We rewrite the proof as follows.

$$\frac{\frac{\mathcal{P}'}{\text{conn}(A, B)} \quad \frac{\mathcal{Q}}{\text{path}(B, C)}}{\text{path}(A, C)} \text{STEP}$$

If \mathcal{Q} does not contain applications of rule **TRANS**, then this is the base case and the thesis follows immediately. Otherwise, the thesis follows by induction hypothesis.

Case Step \mathcal{P} ends with an application of rule **STEP**. Thus, for some \mathcal{P}' and \mathcal{P}'' , we are in the following situation.

$$\frac{\frac{\mathcal{P}'}{\text{conn}(A, B')} \quad \frac{\mathcal{P}''}{\text{path}(B', B)}}{\text{path}(A, B')} \text{STEP} \quad \frac{\mathcal{Q}}{\text{path}(B, C)} \text{TRANS}$$

We rewrite the proof as follows.

$$\frac{\frac{\mathcal{P}'}{\text{conn}(A, B')} \quad \frac{\frac{\mathcal{P}''}{\text{path}(B', B)} \quad \frac{\mathcal{Q}}{\text{path}(B, C)}}{\text{path}(B', C)} \text{TRANS}}{\text{path}(A, C)} \text{STEP}$$

The thesis now follows by induction hypothesis.

Case Trans \mathcal{P} ends with an application of rule **TRANS**. Thus, for some \mathcal{P}' and \mathcal{P}'' , we are in the following situation.

$$\frac{\frac{\mathcal{P}'}{\text{path}(A, B')} \quad \frac{\mathcal{P}''}{\text{path}(B', B)}}{\text{path}(A, B)} \text{TRANS} \quad \frac{\mathcal{Q}}{\text{path}(B, C)} \text{TRANS}$$

From the proof of the left-hand premise and the induction hypothesis, we know that there exists a proof \mathcal{R} that does not contain applications of rule TRANS, ends with $\text{path}(A, B)$, and is not bigger than the left-hand side proof

$$\frac{\frac{\mathcal{P}'}{\text{path}(A, B')}}{\text{path}(A, B)} \quad \frac{\mathcal{P}''}{\text{path}(B', B)} \quad \text{TRANS}.$$

Thus we can build the following smaller derivation, by eliminating at least one application of rule TRANS.

$$\frac{\frac{\mathcal{R}}{\text{path}(A, B)} \quad \frac{\mathcal{Q}}{\text{path}(B, C)}}{\text{path}(A, C)} \quad \text{TRANS}$$

Since the derivation is smaller than the one we started with, we can conclude by invoking the induction hypothesis on this derivation.

□

Exercise 3 (!). Prove that if $\text{path}(A, B)$ is provable in the system in fig. 1.4, then there exists a nonempty sequence of provable propositions $\text{conn}(C_1, C'_1), \dots, \text{conn}(C_n, C'_n)$ for some C_1, \dots, C_n such that $n > 0$, $C_1 = A$ (the sequence starts from A), and $C'_n = B$ (the sequence ends at B).

Prove that if $\text{path}(A, B)$ is provable in the system in fig. 1.1, then there exists a nonempty sequence of provable propositions $\text{conn}(C_1, C'_1), \dots, \text{conn}(C_n, C'_n)$ for some C_1, \dots, C_n such that $n > 0$, $C_1 = A$ (the sequence starts from A), and $C'_n = B$ (the sequence ends at B).

Chapter 2

Simple choreographies

Now that we have familiarised ourselves with formal systems based on inference rules, we can proceed to using them for the study of concurrency. We start in this chapter by building our first, and very simple, choreography model. Our aim is to design the simplest possible choreography language that captures the essence of what a choreography model is and how we can use it.

The cornerstone of our study will be the notion of *process*, an independent agent that can perform local computation and communicate with other agents by means of message passing (Input/Output, or I/O for short). Processes are abstract representations of computer programs executed concurrently, each possessing an independent control state and memory. Essentially, what we are going to do is to use inference systems to model concurrent systems that consist of processes communicating with each other.

Example 1. *As guiding example for this chapter, suppose that we want to define a system that consists of two processes, called **Buyer** and **Seller**. Suppose also that we want these two processes to interact as follows:*

1. *Buyer sends the title of a book she wishes to buy to Seller;*
2. *Seller replies to Buyer with the price of the book.*

The description above is informal. However, it gives us some important indications on how a mathematical formalism for choreographies might look like: our description talks about *multiple* processes and how they interact. We are adopting a global view on all the interactions among the processes that we are interested in. More specifically: each step of our protocol talks about *both* the sender and the receiver of the communication; and we are explicitly ordering communications (as in the “Alice and Bob notation” from the Preface).

$$C ::= p \rightarrow q; C \mid 0$$

Figure 2.1: Simple choreographies, syntax.

2.1 Syntax

Our first choreography model is called simple choreographies. The syntax of simple choreographies is given by the grammar in fig. 2.1, where C ranges over a choreography. Choreographies describe interactions between processes. We refer to processes by using process names, ranged over by p, q .

The syntax of simple choreographies is minimalistic. A choreography can either be a term $p \rightarrow q; C$ or term 0 . Term $p \rightarrow q; C$ is an interaction and reads “process p sends a message to process q ; then, the choreography proceeds as C ”. We always assume that p and q are different in interactions $p \rightarrow q$, i.e., $p \neq q$ (meaning that a process cannot send a message to itself). Term 0 is the terminated choreography (no interactions, or end of program, if you like). Sometimes, we refer to terms as *programs* in the remainder.

Example 2. *The following choreography defines the behaviour that we informally described in example 1.*

$$\text{Buyer} \rightarrow \text{Seller}; \text{Seller} \rightarrow \text{Buyer}; 0$$

Note that what we have is actually a rather coarse abstraction of what we described in example 1, because we are not formalising what is being sent from a process to another. For example, the informal description stated that Buyer sends “the title of a book she wishes to buy” to Seller in the first interaction, but our choreography above does not define this part. It simply states that Buyer sends some unspecified message to Seller, and that Seller replies to Buyer afterwards. We are going to add the possibility to specify the content of messages later on.

2.2 Semantics

Now that we can write simple choreographies, we give them a semantic interpretation. We use a relation to do that. Recall from set theory that a relation \mathcal{R} on two sets S and S' is a subset of the product $S \times S'$ (the set of all pairs of elements from S and S'). Given an element $s \in S$ and an element $s' \in S'$, we write $s\mathcal{R}s'$ for $(s, s') \in \mathcal{R}$.

2.2. Semantics

$$\frac{}{\mathbf{p} \rightarrow \mathbf{q}; C \rightarrow C} \text{ COM}$$

Figure 2.2: Simple choreographies, semantics.

Formally, the semantics for simple choreographies is given in terms of a *reduction relation* \rightarrow , which is defined as the smallest relation satisfying the rule displayed in fig. 2.2. Whenever two choreographies C and C' are related by \rightarrow , written $C \rightarrow C'$, we say that there is a reduction from C to C' . A reduction models executing a step of a choreography.

There is only one rule, called COM, which is an axiom: it always allows us to reduce interactions—if a programmer wishes for an interaction to take place, it always will. In the rule, \mathbf{p} , \mathbf{q} , and C are all schematic variables, on which we impose no conditions. So the rule works for all process names and choreographies. (Recall, however, that we assumed $\mathbf{p} \neq \mathbf{q}$ in communication terms, so this is true also here.)

The sets over which the relation \rightarrow is defined are given implicitly: they are evident from the inference rules that define the relation. Specifically, let *SimpleChor* be the set of all choreography terms in our grammar for simpler choreographies—or, equivalently, the *language* generated by that grammar. There is only one inference rule in fig. 2.2, rule COM. The rule relates any choreography of the form $\mathbf{p} \rightarrow \mathbf{q}; C$ to the choreography C . Thus, we know that $\rightarrow \subseteq \text{SimpleChor} \times \text{SimpleChor}$.

It is important that the set of departure of relation \rightarrow (*SimpleChor*) is the same as its set of destination (*SimpleChor*), i.e., that \rightarrow goes from choreographies to choreographies. This property allows us to define the useful concept of strong reduction chain, which permits to observe multiple steps of execution. The “strong” qualifier indicates that there is at least one step in the chain.

Definition 1 (Strong reduction chain). *We say that there is a strong reduction chain from C_1 to C_n whenever there exists a sequence of choreographies (C_1, \dots, C_n) such that $C_i \rightarrow C_{i+1}$ for each $i \in [1, n - 1]$.*

When there is a strong reduction chain from C_1 to C_n , we write $C_1 \rightarrow^+ C_n$ (when showing the intermediate steps is not necessary, only their existence) or $C_1 \rightarrow \dots \rightarrow C_n$ (when we want to show the intermediate steps).

Example 3. *The program in example 2 has the following strong reduction chain:*

$$\text{Buyer} \rightarrow \text{Seller}; \text{Seller} \rightarrow \text{Buyer}; 0 \quad \rightarrow \quad \text{Seller} \rightarrow \text{Buyer}; 0 \quad \rightarrow \quad 0.$$

So we first reduce an interaction where Buyer sends a message to Seller and then we reduce an interaction where Seller sends a message to Buyer, which is exactly the communication flow that we wanted in example 1.

Another way to define relation \rightarrow^+ is: \rightarrow^+ is the transitive closure of \rightarrow . In the next sections, we will also use the reflexive and transitive closure of \rightarrow , written \rightarrow^* and defined as follows. Whenever $C \rightarrow^* C'$ for some C and C' , we say that there is a weak reduction chain from C to C' .

Definition 2 (Weak reduction chain). *We write $C \rightarrow^* C'$ if either:*

Base case: $C = C'$; or,

Inductive case: *there exists C'' such that $C \rightarrow C''$ and $C'' \rightarrow^* C'$.*

If you are familiar with regular expressions, the use of the symbols $+$ and $*$ should ring a bell: \rightarrow^+ means “one or more reductions” and \rightarrow^* means “zero or more reductions”.

Exercise 4. *Prove that $C \rightarrow^+ C'$ implies $C \rightarrow^* C'$. Prove that the converse does not hold.*

Exercise 5. *Prove the following statements.*

1. *For all $C \neq \mathbf{0}$ (all C that are not $\mathbf{0}$), there exists C' such that $C \rightarrow C'$.*
2. *For all $C \neq \mathbf{0}$, there exists C' such that $C \rightarrow^* C'$.*
3. *For all $C \neq \mathbf{0}$, $C \rightarrow^+ \mathbf{0}$.*
4. *For all C , $C \rightarrow^* \mathbf{0}$.*

Chapter 3

From choreographic programs to process programs

Simple choreographies is like a high-level programming language, providing us with a useful abstraction (the communication term) to talk about what we are interested in. In our case, what we are interested in is defining the interactions that we want to take place among our processes—we want to write protocols. However, high-level programs are not particularly useful if we do not know how they can be implemented in practice. Other high-level languages face the same situation: they offer useful abstractions, like functions and objects, but their programs need to be converted to something that the computer can actually execute, like machine code. Here, we are not interested in reaching the detail of machine code, but in bridging the conceptual difference between choreographies and the typical programs that can be executed in concurrent systems.

In a concurrent and distributed system, each process has its own program that is run by the computer that hosts it. To communicate with each other, processes can then send and receive messages to/from each other. A process does not (necessarily) know what programs the other processes are running, only its own. This is different from choreographies, where the behaviour of multiple processes is defined from a global viewpoint.

Example 4. *To implement our scenario from example 1 following the standard methodology for programming concurrent and distributed systems, we would have to produce two programs: one for process **Buyer** and one for process **Seller**.*

*Informally, a program for **Buyer** could look as follows.*

- *Send the title of the book to buy to **Seller**;*

$$\begin{aligned}
 N &::= p \triangleright B \mid N \mid N \mid \mathbf{0} \\
 B &::= p!; B \mid p?; B \mid \mathbf{0}
 \end{aligned}$$

Figure 3.1: Simple processes, syntax.

- *receive the price of the book from Seller.*

For Seller, we could use the following (abstract) program.

- *Receive the title of the book from Buyer;*
- *send the price of the book to Buyer.*

In a nutshell, choreographies give us a global view on the communications to be enacted by the system; whereas realistic concurrent and distributed systems expect to have a program for each process (we call these *process programs*), based on the local view of that process and using send and receive actions (also called Input/Output, or I/O) to interact with the other processes.

Thus, if we want to make choreographies useful, we need the following two elements.

- A formal model for process programs, or *process model*.
- A way to relate our choreography model (simple choreographies) to the process model or, more concretely, choreography programs to process programs.

3.1 Simple processes

In this section we define a simple process model to describe system implementations. We will use it later to build a notion of *correspondence* between what should happen (given as a choreography) and what the system actually does (given as a term in this process model).

We assume that each process is uniquely identified by its name on the network, and that process names can be used to direct messages from one process to each other (similarly to URLs on the web). based on actors.

3.1. Simple processes

3.1.1 Syntax

The syntax of simple processes is given in fig. 3.1. We model systems as networks, ranged over by N . A network N can be: a single process term $p \triangleright B$, where p is the name of the process and B its behaviour; a parallel composition of two networks N and N' , written $N \mid N'$, which enables the processes in N and N' to communicate; or the empty network 0 . A process behaviour, ranged over by B , can be: a send action $p!; B$, read “send a message to process p and then do B ”; a receive action $p?; B$, read “receive a message from process p and then do B ”; or the terminated behaviour 0 .

Example 5. *The two process programs informally described in example 4 can be formalised as follows.*

$$\text{Buyer} \triangleright \text{Seller}!; \text{Seller}?; 0 \mid \text{Seller} \triangleright \text{Buyer}?; \text{Buyer}!; 0$$

3.1.2 Semantics: discussion

Similarly to what we have done for simple choreographies, we can equip simple processes with a semantics based on reductions of the form $N \rightarrow N'$. The idea is that each reduction should model a step in the execution of a network. However, this requires a few extra ingredients compared to the semantics of simple choreographies, so we make a few considerations before giving the general definition.

The essential rule defining reductions for simple processes is the following axiom, rule COM.

$$\frac{}{p \triangleright q!; B \mid q \triangleright p?; B' \rightarrow p \triangleright B \mid q \triangleright B'} \text{COM}$$

Rule COM models communications, by matching a send action by process p towards process q with a receive action at q waiting for a message from p . Each process then proceeds with its respective continuation (B for p , B' for q).

Rule COM is simple, because it focuses only on the components that it needs (the sender and the receiver). However, if we designed a semantics that includes only that rule, the result would be too limited. For example, consider the network from example 5:

$$\text{Buyer} \triangleright \text{Seller}!; \text{Seller}?; 0 \mid \text{Seller} \triangleright \text{Buyer}?; \text{Buyer}!; 0.$$

By rule COM, we have the following reduction:

$$\begin{aligned} &\text{Buyer} \triangleright \text{Seller}!; \text{Seller}?; 0 \mid \text{Seller} \triangleright \text{Buyer}?; \text{Buyer}!; 0 \\ &\rightarrow \\ &\text{Buyer} \triangleright \text{Seller}?; 0 \mid \text{Seller} \triangleright \text{Buyer}!; 0. \end{aligned}$$

Chapter 3. From choreographic programs to process programs

So far so good. But what about the reductum (the term of the reduction)? Intuitively, we would expect to be able to reduce the reductum

$$\text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \mid \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \quad (3.1)$$

as follows, because **Buyer** and **Seller** are performing compatible actions—**Buyer** wants to receive from **Seller**, and **Seller** wants to send to **Buyer**.

$$\text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \mid \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \rightarrow \text{Buyer} \triangleright \mathbf{0} \mid \text{Seller} \triangleright \mathbf{0}$$

Sadly, this does not work, because rule **COM** requires the sender to appear on the left-hand side of the parallel composition.

One solution could be to have another rule, **MOC** (the reverse of **COM**), that allows us to reduce networks where the sender appears on the right-hand side.

$$\frac{}{p \triangleright q?; B \mid q \triangleright p!; B' \rightarrow p \triangleright B \mid q \triangleright B'} \text{MOC}$$

While **MOC** might look like a good idea at first, we quickly run into trouble again with other networks. What if we have a terminated process **Bystander** in between **Buyer** and **Seller** in our example?

$$\text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \mid \text{Bystander} \triangleright \mathbf{0} \mid \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \quad (3.2)$$

We cannot apply rule **MOC** to the network above, because the rule does not allow for any terms to appear in between the two interacting processes.

Instead of designing an ad-hoc rule for each corner case, we adopt the standard technique of defining a *structural relation* [Sangiorgi and Walker, 2001]. A structural relation relates terms that are intuitively equivalent and should behave in the same way, but are not syntactically equal. We will denote our structural relation for processes with \preceq .

To gain some intuition on what terms \preceq should relate, consider again the network with the bystander in eq. (3.2). Since process **Bystander** is not performing any actions at all, we should be able to disregard it entirely when reasoning about our network. In other words, we wish for the following to hold.

$$\begin{array}{l} \text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \\ \mid \text{Bystander} \triangleright \mathbf{0} \\ \mid \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \end{array} \preceq \begin{array}{l} \text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \\ \mid \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \end{array}$$

Observe that, in the right-hand side, we ended up in the same situation as in eq. (3.1): we cannot apply rule **COM** because the sender and the receiver do not appear in the right order. However, it does not make sense that order matters in parallel compositions. If two processes wish to speak to

3.1. Simple processes

each other, then they should just be able to do so. Why should the order in which they are “placed in the network” matter? Formally, we wish for the following.

$$\text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \mid \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \preceq \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \mid \text{Buyer} \triangleright \text{Seller?}; \mathbf{0}$$

Another wish is that \preceq is not finer than syntactic equality, i.e., if two terms are exactly the same, then they are also structurally related. Formally, for any two networks N and N' , $N = N'$ implies that $N \preceq N'$.

If we had a relation \preceq that grants our wishes (and we are going to define it), then we could use it in our reduction semantics by adopting the following rule.

$$\frac{N \preceq N_1 \quad N_1 \rightarrow N_2 \quad N_2 \preceq N'}{N \rightarrow N'} \text{STRUCT}$$

Rule STRUCT closes \rightarrow under \preceq . With this rule, we can finally derive that our network with the bystander from eq. (3.2) can reduce and terminate.

$$\frac{\begin{array}{c} \text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \\ \mid \text{Bystander} \triangleright \mathbf{0} \\ \mid \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \end{array} \preceq \begin{array}{c} \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \\ \mid \text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \end{array} \quad \begin{array}{c} \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \\ \mid \text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \end{array} \rightarrow \begin{array}{c} \text{Seller} \triangleright \mathbf{0} \\ \mid \text{Buyer} \triangleright \mathbf{0} \end{array} \quad \begin{array}{c} \text{Seller} \triangleright \mathbf{0} \\ \mid \text{Buyer} \triangleright \mathbf{0} \end{array} \rightarrow \begin{array}{c} \text{Buyer} \triangleright \mathbf{0} \\ \mid \text{Seller} \triangleright \mathbf{0} \end{array}}{\begin{array}{c} \text{Buyer} \triangleright \text{Seller?}; \mathbf{0} \\ \mid \text{Bystander} \triangleright \mathbf{0} \\ \mid \text{Seller} \triangleright \text{Buyer!}; \mathbf{0} \end{array} \rightarrow \begin{array}{c} \text{Buyer} \triangleright \mathbf{0} \\ \mid \text{Seller} \triangleright \mathbf{0} \end{array}} \text{STRUCT}$$

3.1.3 Semantics: definition

We move to the formal definition of the semantics of simple processes. Since the definition of the reduction relation \rightarrow depends on the definition of the structural relation \preceq , we define the latter first.

Structural relation

First, a useful abbreviation: we will write $N \equiv N'$ whenever $N \preceq N'$ and $N' \preceq N$ (the symmetric cases).

We define the structural relation \preceq as the smallest relation that satisfies the axioms displayed in fig. 3.2.

Rules PC and PA capture that parallel composition is commutative (PC) and associative (PA), so that we can disregard the order in which networks are composed. Rule GCN garbage collects a terminated network, and rule GCP transforms a terminated process into a terminated network (which can be later collected by rule GCN).

$$\begin{array}{c}
 \overline{N_1 \mid N_2 \equiv N_2 \mid N_1} \text{ PC} \quad \overline{(N_1 \mid N_2) \mid N_3 \equiv N_1 \mid (N_2 \mid N_3)} \text{ PA} \\
 \overline{N \mid \mathbf{0} \preceq N} \text{ GCN} \quad \overline{p \triangleright \mathbf{0} \preceq \mathbf{0}} \text{ GCP} \\
 \overline{N \preceq N} \text{ REFL} \quad \frac{N \preceq N' \quad N' \preceq N''}{N \preceq N''} \text{ TRANS} \quad \frac{N \preceq N'}{\mathcal{C}[N] \preceq \mathcal{C}[N']} \text{ CTX}
 \end{array}$$

Figure 3.2: Simple processes, structural relation.

Rule REFL makes the structural relation reflexive, which makes it include syntactic equality (a network is always structurally related to itself).

Rule TRANS makes the structural relation transitive, so that we can “use it” multiple times, e.g., to permute the processes in a parallel composition until we reach the configuration that we desire.

Rule CTX closes the structural relation up to any context \mathcal{C} , which we formally define now.

Definition 3 (Context). *A context, denoted \mathcal{C} , is obtained when the hole \bullet (a special terminal symbol) replaces one occurrence of $\mathbf{0}$ in a network.*

Given a context \mathcal{C} and a network N , then $\mathcal{C}[N]$ is the network obtained by replacing the hole \bullet in \mathcal{C} with N .

Example 6. *Let the context \mathcal{C} be defined as follows. We put explicit parentheses for grammatical clarity.*

$$\mathcal{C} \triangleq (\text{Buyer} \triangleright \text{Seller}?; \mathbf{0} \mid \bullet) \mid \text{Seller} \triangleright \text{Buyer}!; \mathbf{0}$$

Then, the network with the bystander in eq. (3.2) can be obtained by replacing the hole in the context \mathcal{C} with the bystander process:

$$\mathcal{C}[\text{Bystander} \triangleright \mathbf{0}] = (\text{Buyer} \triangleright \text{Seller}?; \mathbf{0} \mid \text{Bystander} \triangleright \mathbf{0}) \mid \text{Seller} \triangleright \text{Buyer}!; \mathbf{0}.$$

Example 7. *Thanks to contexts, rule CTX allows us to apply the structural relation to sub-terms of networks.*

Let \mathcal{C} be defined as in example 6 and \mathcal{D} be defined as follows.

$$\mathcal{D} \triangleq \bullet \mid \text{Seller} \triangleright \text{Buyer}!; \mathbf{0}$$

3.1. Simple processes

$$\begin{array}{c}
\frac{}{p \triangleright q!; B \mid q \triangleright p?; B'} \text{ COM} \\
\\
\frac{N_1 \rightarrow N'_1}{N_1 \mid N_2 \rightarrow N'_1 \mid N_2} \text{ PAR} \quad \frac{N \preceq N_1 \quad N_1 \rightarrow N_2 \quad N_2 \preceq N'}{N \rightarrow N'} \text{ STRUCT}
\end{array}$$

Figure 3.3: Simple processes, semantics.

Notice that $\mathcal{C}[0] = \mathcal{D}[\text{Buyer} \triangleright \text{Seller?}; 0 \mid 0]$. Then, we can derive:

$$\frac{\frac{\frac{}{\text{Bystander} \triangleright 0 \preceq 0} \text{ GCP}}{\mathcal{C}[\text{Bystander} \triangleright 0] \preceq \mathcal{C}[0]} \text{ CTX} \quad \frac{\frac{}{\text{Buyer} \triangleright \text{Seller?}; 0 \mid 0 \preceq \text{Buyer} \triangleright \text{Seller?}; 0} \text{ GCN}}{\mathcal{D}[\text{Buyer} \triangleright \text{Seller?}; 0 \mid 0] \preceq \mathcal{D}[\text{Buyer} \triangleright \text{Seller?}; 0]} \text{ CTX}}{\frac{(\text{Buyer} \triangleright \text{Seller?}; 0 \mid \text{Bystander} \triangleright 0) \preceq \text{Buyer} \triangleright \text{Seller?}; 0 \mid \text{Seller} \triangleright \text{Buyer!}; 0} \text{ TRANS}}$$

In the remainder, we will not show explicit derivations for the structural relation. It should be evident, for example, that from the conclusion that we have just derived, we can obtain the following by applying rules TRANS (transitivity) and PC (commutativity of parallel).

$$\frac{(\text{Buyer} \triangleright \text{Seller?}; 0 \mid \text{Bystander} \triangleright 0) \preceq \text{Buyer} \triangleright \text{Seller?}; 0 \mid \text{Seller} \triangleright \text{Buyer!}; 0}{\text{Seller} \triangleright \text{Buyer!}; 0 \mid \text{Buyer} \triangleright \text{Seller?}; 0}$$

Reductions

The semantics of simple processes is given by the reduction rules displayed in fig. 3.3.

Rule COM models communications, by matching a send action by process p towards process q with a receive action at q waiting for a message from p . Each process then proceeds with its respective continuation (B for p , B' for q).

Rule PAR allows for reductions to happen inside of a sub-network.

Rule STRUCT allows us to apply the structural relation \preceq when deriving reductions.

Rule PAR allows for reductions to happen in a sub-network.

Rule STRUCT closes reductions under the structural relation \preceq , allowing us to disregard structural differences in networks.

Exercise 6. Prove that the following reductions hold.

$$\begin{aligned} & \text{Buyer} \triangleright \text{Seller!}; \text{Seller?}; 0 \mid \text{Seller} \triangleright \text{Buyer?}; \text{Buyer!}; 0 \\ & \rightarrow \\ & \text{Buyer} \triangleright \text{Seller?}; 0 \mid \text{Seller} \triangleright \text{Buyer!}; 0 \end{aligned}$$

$$\text{Buyer} \triangleright \text{Seller?}; 0 \mid \text{Seller} \triangleright \text{Buyer!}; 0 \rightarrow 0$$

$$(\text{Buyer} \triangleright \text{Seller?}; 0 \mid \text{Bystander} \triangleright 0) \mid \text{Seller} \triangleright \text{Buyer!}; 0 \rightarrow 0$$

Exercise 7 (!). *Prove the following statement.*

If $N_1 \preceq N_2$ and $N_1 \preceq N_3$, then there exists N' such that $N_2 \preceq N'$ and $N_3 \preceq N'$.

3.2 Endpoint Projection (EPP)

The network in example 5 works as intended, but we had to come up with it manually. If we could figure out a mechanical method of going from a choreography (which formalises what we want) to a network (which formalises an implementation), we would save time. If we could also prove that such method *always gives us a correct result*, we would also save ourselves the potential mistakes that come from the manual activity of writing a process network that should implement what we want.

3.2.1 From choreographies to processes

To gain some intuition on how we could develop the method we want, we can look at our examples. Let us see our choreography from example 2 again:

$$\text{Buyer} \rightarrow \text{Seller}; \text{Seller} \rightarrow \text{Buyer}; 0.$$

It is evident, albeit informally, that our network from example 5 implements exactly the interactions defined in the choreography:

$$\text{Buyer} \triangleright \text{Seller!}; \text{Seller?}; 0 \mid \text{Seller} \triangleright \text{Buyer?}; \text{Buyer!}; 0.$$

This informal correspondence is preserved by reductions—remember, reductions model execution, if you think in computational terms. Indeed, whenever we take a step in the reduction chain shown in example 3 (for the choreography), we can “mimic” it by following the reduction chain of example 5 (for the process network), and vice versa (if we take a step for the process network, we can mimic it for the choreography).

3.2. Endpoint Projection (EPP)

The intuition that we can gain from our examples is that a network “implements” a choreography if the actions performed by processes give rise to the interactions described in the choreography. Therefore, an automatic method that produces networks from choreographies should generate a network that consists of the processes described in the choreography, and the behaviour of each of these processes should be the actions that the process needs to perform to implement the interactions that it is involved in in the choreography.

We can now move to formally defining our desired method, as a function from choreographies to networks. This function is commonly called *EndPoint Projection* (EPP for short), since it projects each interaction in the choreography to the local action that each process (an endpoint) should perform in the network [Qiu et al., 2007, Lanese et al., 2008, Carbone et al., 2012]. Indeed, we can think of an interaction like $\text{Buyer} \rightarrow \text{Seller}$ as consisting of two parts, i.e., the send action by Buyer and the receive action by Seller . So the send action that the process implementing Buyer should perform is the first component of the interaction, and the receive action by Seller is the second component.

Let $\text{procs}(C)$ be the set of process names used in C . We can define this function inductively on the structure of C , as follows.

$$\begin{aligned}\text{procs}(\mathfrak{p} \rightarrow \mathfrak{q}; C) &= \{\mathfrak{p}, \mathfrak{q}\} \cup \text{procs}(C) \\ \text{procs}(\mathbf{0}) &= \emptyset\end{aligned}$$

We write $\llbracket C \rrbracket$ for the EPP of a choreography C .

Definition 4 (EndPoint Projection (EPP)). *The EPP of a choreography C , denoted $\llbracket C \rrbracket$, is defined as:*

$$\llbracket C \rrbracket = \prod_{\mathfrak{p} \in \text{procs}(C)} \mathfrak{p} \triangleright \llbracket C \rrbracket_{\mathfrak{p}}.$$

In definition 5, the notation $\prod_{\mathfrak{p} \in \text{procs}(C)} \mathfrak{p} \triangleright \llbracket C \rrbracket_{\mathfrak{p}}$ stands for “the parallel composition of all $\mathfrak{p} \triangleright \llbracket C \rrbracket_{\mathfrak{p}}$ such that \mathfrak{p} is in $\text{procs}(C)$ ”. The auxiliary function $\llbracket C \rrbracket_{\mathfrak{p}}$ —not to be confused with $\llbracket C \rrbracket$ —projects the actions that process \mathfrak{p} should perform in order to implement its part in choreography C . We call $\llbracket C \rrbracket_{\mathfrak{p}}$ a behaviour projection, since it outputs a behaviour B . It is inductively defined by the rules given in fig. 3.4.

Example 8. Let $C_{\text{BuyerSeller}}$ be the choreography in example 2. We recall it here (\triangleq stands for “defined as”):

$$C_{\text{BuyerSeller}} \triangleq \text{Buyer} \rightarrow \text{Seller}; \text{Seller} \rightarrow \text{Buyer}; \mathbf{0}.$$

$$\llbracket p \rightarrow q; C \rrbracket_r = \begin{cases} q!; \llbracket C \rrbracket_r & \text{if } r = p \\ p?; \llbracket C \rrbracket_r & \text{if } r = q \\ \llbracket C \rrbracket_r & \text{otherwise} \end{cases}$$

$$\llbracket 0 \rrbracket_p = 0$$

Figure 3.4: Behaviour projection for simple choreographies.

The process names in $C_{\text{BuyerSeller}}$ are:

$$\text{procs}(C_{\text{BuyerSeller}}) = \{\text{Buyer}, \text{Seller}\}.$$

The behaviour projection for Buyer is:

$$\llbracket C_{\text{BuyerSeller}} \rrbracket_{\text{Buyer}} = \text{Seller!}; \text{Seller?}; 0.$$

The EPP of $C_{\text{BuyerSeller}}$ is exactly the network that we defined in example 5:

$$\llbracket C_{\text{BuyerSeller}} \rrbracket = \text{Buyer} \triangleright \text{Seller!}; \text{Seller?}; 0 \mid \text{Seller} \triangleright \text{Buyer?}; \text{Buyer!}; 0.$$

Exercise 8. Write the outputs of procs and $\llbracket \cdot \rrbracket$ (EPP) for the choreography

$$p \rightarrow q; r \rightarrow q; r \rightarrow s; q \rightarrow p; 0.$$

Exercise 9. What are the reduction chains for the choreography in exercise 8 and (the network resulting from) its EPP? Do you think that they informally correspond to one another? (We have not formally defined correspondence yet.)

Exercise 10. Is the following statement true?

Let $N = \llbracket C \rrbracket$. If $N \rightarrow N'$ for some N' , then there exists C' such that $C \rightarrow C'$ and $N' \preceq \llbracket C' \rrbracket$.

3.2.2 Towards a correct EPP

Unfortunately, the statement in exercise 10 does not hold. More specifically, the framework of simple choreographies does not support a key desirable property: the EPP of a choreography should only do what the original choreography prescribes.

The counterexample is simple. Take the following choreography:

$$C_{\text{problem}} \triangleq p \rightarrow q; r \rightarrow s; 0. \tag{3.3}$$

3.2. Endpoint Projection (EPP)

The EPP of this choreography is:

$$\llbracket C_{\text{problem}} \rrbracket = p \triangleright q!; 0 \mid q \triangleright p?; 0 \mid r \triangleright s!; 0 \mid s \triangleright r?; 0.$$

We have the following reduction for this network, by synchronising r with s :

$$p \triangleright q!; 0 \mid q \triangleright p?; 0 \mid r \triangleright s!; 0 \mid s \triangleright r?; 0 \rightarrow p \triangleright q!; 0 \mid q \triangleright p?; 0.$$

However, this reduction cannot be mimicked by C_{problem} , which can only reduce the first interaction between p and q according to the semantics given in the previous lecture notes.

The example above shows that our framework is not sound yet, because the projection of a choreography can perform “extra” reductions with respect to the choreography. There are two ways to fix this.

Forbidding independent sequences of interactions One way is to say that choreographies like C_{problem} are “forbidden”, because the sequence of interactions $p \rightarrow q; r \rightarrow s$ is not enforced by any causality relation between the two interactions. More specifically, since the processes p , q , r , and s are all different, they operate independently—as the reduction for the projection of C_{problem} shows. However, if the process names of the two interactions intersected in any way, this problem would not appear. For example, consider the choreography $p \rightarrow q; r \rightarrow q; 0$. Its projection reduces as expected by the choreography because process q necessarily needs to complete the first interaction before participating in the second.

Exercise 11. *Check that the EPP of $p \rightarrow q; r \rightarrow q; 0$ reduces as expected by the choreography (i.e., their respective reduction chains match).*

Detecting sequences of interactions that do not have causal dependencies can be done mechanically. Thus, it is possible to automatically detect whether a choreography respects the condition of not having sequences of independent interactions, as in C_{problem} . Forbidding programmers to write choreographies like C_{problem} was a popular approach (and still is in some works, when useful) in the first studies on choreographies, like [Fu et al., 2005, Qiu et al., 2007, Carbone et al., 2007].

Out of order execution The other way to solve our issue with sequences of independent interactions is to extend the semantics of choreographies to correctly capture the parallel nature of processes. Going back to C_{problem} again, if we could somehow design a semantics that allowed for the following reduction

$$p \rightarrow q; r \rightarrow s; 0 \rightarrow p \rightarrow q; 0$$

$$C ::= p \rightarrow q; C \mid 0$$

Figure 3.5: Simple concurrent choreographies, syntax.

$$\frac{}{p \rightarrow q; C \rightarrow C} \text{COM} \quad \frac{C \preceq C_1 \quad C_1 \rightarrow C_2 \quad C_2 \preceq C'}{C \rightarrow C'} \text{STRUCT}$$

Figure 3.6: Simple concurrent choreographies, semantics.

then we would be fine, because that is the reduction that we need to match the problematic one that the EPP of the choreography can do. Observe that this reduction does not respect the order in which instructions are given in the choreography. This kind of semantics is typically called “out-of-order execution”. The idea is widespread in many domains. For example, modern CPUs and/or language runtimes may change the order in which instructions are executed to increase performance, when it is safe to do so—a typical example is the single-threaded imperative code `x++; y++;`, where the order in which the two increments are done is influential and the runtime may thus decide to parallelise it.

Since the inception of out-of-order execution for choreographies [Carbone and Montesi, 2013], similar ideas have been adopted in different works [Deniélou and Yoshida, 2013, Honda et al., 2016]. In the next section, where we develop the technical details, we borrow the formulation by Cruz-Filipe and Montesi [2016].

3.2.3 Simple Concurrent Choreographies

We update our model of simple choreographies to capture concurrent execution of different processes.

The syntax of choreographies remains unchanged. It is displayed in fig. 3.5.

The semantics of choreographies, instead, needs some updating to capture out-of-order execution of interactions. We obtain this by adding a structural relation \preceq for choreographies¹. The reduction rules are given in fig. 3.6. The only change is the addition of rule STRUCT, which closes reductions under \preceq .

¹We use the same symbol as for the structural relation for networks, since we can easily distinguish them from the context (they operate on different domains, one choreographies and the other networks).

3.2. Endpoint Projection (EPP)

$$\frac{\{p, q\} \# \{r, s\}}{p \rightarrow q; r \rightarrow s \equiv r \rightarrow s; p \rightarrow q} \text{SWAP}$$

Figure 3.7: Simple concurrent choreographies, structural precongruence.

We define \preceq as the smallest relation that satisfies the rule displayed in fig. 3.7 and is reflexive, transitive, and closed under (choreographic) contexts. We do not give explicit inference rules for the last three properties (reflexive, transitive, context closure), as these should be intuitive by now. Recall that $C \equiv C'$ stands for $C \preceq C'$ and $C' \preceq C$. (Later on, we will see rules that use \preceq in only one direction also for choreographies.) Rule SWAP states that two interactions $p \rightarrow q$ and $r \rightarrow s$ can be exchanged in a choreography if the processes p, q, r , and s are distinct (they are all different). This is captured by the premise $\{p, q\} \# \{r, s\}$, which states that the sets $\{p, q\}$ and $\{r, s\}$ must be disjoint. Formally, given two sets S and S' , $S \# S'$ is a shortcut notation for $S \cap S' = \emptyset$ (empty intersection). Notice that we slightly abuse notation in rule SWAP: a term $p \rightarrow q; r \rightarrow s$ is only “partial”, in the sense that it is not a valid complete term for the grammar of choreographies, it can only be part of a bigger complete term. Context closure of \preceq means that we can apply that rule to all subterms of a choreography.

Example 9. *This is the reduction we wished we could do in section 3.2.2.*

$$p \rightarrow q; r \rightarrow s; 0 \rightarrow p \rightarrow q; 0.$$

We can now perform it with our new semantics. Here is the derivation:

$$\frac{\begin{array}{c} p \rightarrow q; \quad r \rightarrow s; \\ r \rightarrow s; 0 \preceq p \rightarrow q; 0 \end{array} \quad \frac{\text{COM}}{p \rightarrow q; 0 \preceq p \rightarrow q; 0} \quad \frac{\text{STRUCT}}{p \rightarrow q; r \rightarrow s; 0 \rightarrow p \rightarrow q; 0}$$

Exercise 12. *Show all the possible reduction chains of the following choreography.*

$$p \rightarrow q; r \rightarrow s; q \rightarrow r; 0$$

Exercise 13. *Prove the following statement.*

Let $\llbracket C \rrbracket = N$. If $C \preceq C'$ for some C' , then $N \preceq \llbracket C' \rrbracket$.

We can now formally state that EPP is correct, in the sense that the behaviour implemented by the network projected from a choreography is exactly the one defined in the choreography.

Chapter 3. From choreographic programs to process programs

Theorem 2 (Operational Correspondence). *Let $\llbracket C \rrbracket = N$. Then,*

Completeness *If $C \rightarrow C'$ for some C' , then there exists N' such that $N \rightarrow N'$ and $N' \preceq \llbracket C' \rrbracket$.*

Soundness *If $N \rightarrow N'$ for some N' , then there exists C' such that $C \rightarrow C'$ and $N' \preceq \llbracket C' \rrbracket$.*

Intuitively, the completeness part means that the network generated by the EPP of a choreography does *all* that the choreography says. Conversely, the soundness part means that the network generated by the EPP of a choreography does *only* what the choreography says. The two parts combined give us correctness: the network generated by EPP does exactly what is defined in the originating choreography. So what happens is what we want, finally!

The correctness of EPP gives us a powerful result for free: the EPP of a choreography never gets stuck (for example, there cannot be deadlocks). Intuitively, this works because interactions in choreographies are written atomically, in the sense that they specify both the send and receive actions needed for the communication in a single term. To have a deadlock, one typically needs a language where send and receive actions are defined separately (hence the opportunity for mistakes).

Let us formalise this result. First, we observe that choreographies never get stuck.

Theorem 3 (Progress for choreographies). *Let C be a choreography. Then, either $C = \mathbf{0}$ (C has terminated) or there exists C' such that $C \rightarrow C'$.*

Proof. By cases on C . If $C = \mathbf{0}$, the thesis follows immediately. Otherwise, we can just apply rule COM. \square

We now combine theorem 3 with the completeness part of theorem 2, which respectively say that “a choreography can always reduce until it terminates” and “the EPP of a choreography can always do what the choreography does”. This means that “the EPP of a choreography can always reduce until it terminates”, as formalised below.

Corollary 1 (Progress for EPP). *Let $\llbracket C \rrbracket = N$. Then, either $N = \mathbf{0}$ (N has terminated) or there exists N' such that $N \rightarrow N'$.*

Proof. Direct consequence of theorems 2 and 3. \square

Chapter 4

Statefulness

4.1 Local Computation

Now that we know the basic soundness principles of choreographies and EPP, we can play with extending our model such that we can capture more interesting—and realistic—examples.

In this section, we equip processes with the ability to perform local computation. This will enable us to capture our initial scenario from example 1 more precisely, i.e., we want to define the *content* of the messages exchanged by Buyer and Seller.

4.1.1 Stateful Choreographies

Syntax We augment the syntax of choreographies with *local functions*—ranged over by f, g, \dots —which can be used by processes to perform local computation. The new syntax is given in fig. 4.1.

The key idea is that now processes are equipped with their own local memories, which they can manipulate through computation. The new term $\mathbf{p}.f$ reads “process \mathbf{p} stores the result of function f in its memory”. The new communication term $\mathbf{p}.f \rightarrow \mathbf{q}.g; C$ reads “process \mathbf{p} sends the result of computing function f to \mathbf{q} ; \mathbf{q} then computes function g according to the

$$\begin{aligned} C &::= \eta; C \mid \mathbf{0} \\ \eta &::= \mathbf{p}.f \rightarrow \mathbf{q}.g \mid \mathbf{p}.f \end{aligned}$$

Figure 4.1: Stateful choreographies, syntax.

$$\begin{array}{c}
 \frac{f(\sigma(\mathbf{p})) \downarrow v}{\langle \mathbf{p}.f; C, \sigma \rangle \rightarrow \langle C, \sigma[\mathbf{p} \mapsto v] \rangle} \text{LOCAL} \\
 \\
 \frac{f(\sigma(\mathbf{p})) \downarrow v \quad g(\sigma(\mathbf{q}), v) \downarrow u}{\langle \mathbf{p}.f \rightarrow \mathbf{q}.g; C, \sigma \rangle \rightarrow \langle C, \sigma[\mathbf{q} \mapsto u] \rangle} \text{COM} \\
 \\
 \frac{C \preceq C_1 \quad \langle C_1, \sigma \rangle \rightarrow \langle C_2, \sigma' \rangle \quad C_2 \preceq C'}{\langle C, \sigma \rangle \rightarrow \langle C', \sigma' \rangle} \text{STRUCT}
 \end{array}$$

Figure 4.2: Stateful choreographies, semantics.

received message and its local memory, and updates its memory with the result”.

Notice that we now use η to range over choreographic statements, i.e., communications $(\mathbf{p}.f \rightarrow \mathbf{q}.g)$ and internal computation steps $(\mathbf{p}.f)$. This is convenient for reasoning about process names in statements. We update the definition of **procs** as follows.

$$\begin{aligned}
 \text{procs}(\eta; C) &= \text{procs}(\eta) \cup \text{procs}(C) \\
 \text{procs}(\mathbf{0}) &= \emptyset \\
 \text{procs}(\mathbf{p}.f \rightarrow \mathbf{q}.g) &= \{\mathbf{p}, \mathbf{q}\} \\
 \text{procs}(\mathbf{p}.f) &= \{\mathbf{p}\}
 \end{aligned}$$

Semantics Now that processes have functions that may refer to their memories, the execution of a choreography depends on the state of process memories. To formalise this, we need to formulate reductions on more than just choreographies: we will generalise reductions to consider pairs of choreographies and memory states.

It is convenient to abstract from how process memory is concretely implemented, since different processes may use different kinds of data structures. Let v, u, \dots range over memory states (or values), which we leave unspecified. Also, let σ range over global memory states, mapping process names to values. Intuitively, σ maps each process to its memory state. For example, $\sigma(\mathbf{p}) = v$ means that process \mathbf{p} has v as memory. We assume that σ s are total functions (they are never undefined).

We define a semantics for stateful choreographies in terms of reductions $\langle C, \sigma \rangle \rightarrow \langle C', \sigma' \rangle$, where $\langle C, \sigma \rangle$ is a *runtime configuration*. The rules defining \rightarrow are given in fig. 4.2. It is based as usual on a structural relation \preceq , defined by the rule in fig. 4.3.

4.1. Local Computation

$$\frac{\text{procs}(\eta) \# \text{procs}(\eta')}{\eta; \eta' \equiv \eta'; \eta} \text{SWAP}$$

Figure 4.3: Stateful choreographies, structural precongruence.

Let us look at rule LOCAL first, which gives a semantics for local computations. The premise uses the evaluation operator \downarrow . We write $f(v_1, \dots, v_n) \downarrow u$ when the evaluation of function f —where its parameters are instantiated with the values v_1, \dots, v_n —returns the value u . We do not define how \downarrow works concretely, since that depends on the language in which functions are written, which is an implementation detail from our perspective here. In rule LOCAL, we pass the current memory state of \mathbf{p} — $\sigma(\mathbf{p})$ —to f and get a value v , which then becomes the new state for process \mathbf{p} . The notation $\sigma[\mathbf{p} \mapsto v]$ is a mapping update and means “ σ , but where \mathbf{p} is now mapped to v ”. Formally:

$$(\sigma[\mathbf{p} \mapsto v])(\mathbf{q}) = \begin{cases} v & \text{if } \mathbf{q} = \mathbf{p} \\ \sigma(\mathbf{q}) & \text{otherwise} \end{cases}.$$

Rule COM is the natural extension of interactions to include internal computation. In the first premise, we evaluate the function f used by the sender \mathbf{p} under its memory state, getting a value v . Then, we evaluate the function g used by the receiver under the memory state of the receiver and the value v (the message received from \mathbf{p}), obtaining a value u . The memory of the receiver \mathbf{q} is then updated to become this value.

We assume that evaluating a local function always terminates (it never takes infinite time). In practice, this means that local computation may yield error values, and that infinite computations may be interrupted by timeouts. We abstract from such details, since what we are interested in here is communications, not the algorithmic details of internal computation. It is easy to plug in existing techniques for ensuring that the parameters passed to local functions are always of the right type, as shown in [Cruz-Filipe and Montesi, 2017].

Rule SWAP is updated to deal with all kinds of choreographic statements, be they interactions or internal computations.

Modelling variables Mainstream programming languages typically allow programmers to manipulate different *variables* whose values reside in memory. Luckily, we do not need to extend our model to capture this style, since we can already capture it by using functions and a few conventions.

Let x, y, z, \dots range over variable names (variables for short). A variable mapping h is a total function that maps variables to values. From now on,

we adopt the following shortcut notation¹: $p.f \rightarrow q.x$ stands for $p.f \rightarrow q.set^x$, where set^x is the function that replaces x in the variable mapping of q with the value received from p . Formally, given x , the evaluation of set^x is defined as:

$$set^x(h, v) \downarrow h[x \mapsto v].$$

Exercise 14. *Prove the following statement.*

For all $\langle p.f \rightarrow q.x; C, \sigma \rangle$ such that $\sigma(q) = h$ and h is a variable mapping, we have that

$$\langle p.f \rightarrow q.x; C, \sigma \rangle \rightarrow \langle C, \sigma[q \mapsto h'] \rangle$$

where $f(\sigma(p)) = v$ and $h' = h[x \mapsto v]$.

Conversely, it is useful to have a shortcut notation for *sending* the content of a variable. Thus, from now on, we adopt also another shortcut notation: $p.x \rightarrow q.g$ stands for $p.get^x \rightarrow q.g$, where get^x is the function that returns the value of variable x from the variable mapping of p . Formally, given x , the evaluation of get^x is defined as:

$$get^x(h) \downarrow h(x).$$

Example 10. *We can finally give a precise choreography for our example introduced in Part 1, including computation and message contents as well. Recall the informal description of the example:*

1. Buyer sends the title of a book she wishes to buy to Seller;
2. Seller replies to Buyer with the price of the book.

A corresponding choreography that defines this behaviour is

$$\text{Buyer.title} \rightarrow \text{Seller.x}; \text{Seller.cat}(x) \rightarrow \text{Buyer.price}; 0$$

where title is a variable, cat is a function (cat stands for catalogue, if you like) that given a book title returns the price for it, and price is a variable.

Exercise 15. *Let σ be such that $\sigma(p) = h_p$ and $\sigma(q) = h_q$ for some variable mappings h_p and h_q . Also, let $h_p(\text{title}) = \text{"Flowers for Algernon"}$ and cat be a function such that $\text{cat}(\text{"Flowers for Algernon"}) = 100$. Show the reduction chain for the choreography in the previous example:*

$$\text{Buyer.title} \rightarrow \text{Seller.x}; \text{Seller.cat}(x) \rightarrow \text{Buyer.price}; 0.$$

¹Shortcut notations for programming languages are sometimes called syntactic sugar.

4.1. Local Computation

$$\begin{aligned}
N &::= \mathbf{p} \triangleright_v B \mid N \mid N \mid \mathbf{0} \\
B &::= \mathbf{p}!f; B \mid \mathbf{p}?f; B \mid f; B \mid \mathbf{0}
\end{aligned}$$

Figure 4.4: Stateful processes, syntax.

$$\begin{aligned}
&\frac{f(v) \downarrow v' \quad g(u, v') \downarrow u'}{\mathbf{p} \triangleright_v \mathbf{q}!f; B \mid \mathbf{q} \triangleright_u \mathbf{p}?g; B' \rightarrow \mathbf{p} \triangleright_v B \mid \mathbf{q} \triangleright_{u'} B'} \text{COM} \\
&\frac{f(v) \downarrow u}{\mathbf{p} \triangleright_v f; B \rightarrow \mathbf{p} \triangleright_u B} \text{LOCAL} \quad \frac{N_1 \rightarrow N'_1}{N_1 \mid N_2 \rightarrow N'_1 \mid N_2} \text{PAR} \\
&\frac{N \preceq N_1 \quad N_1 \rightarrow N_2 \quad N_2 \preceq N'}{N \rightarrow N'} \text{STRUCT}
\end{aligned}$$

Figure 4.5: Stateful processes, semantics.

4.1.2 Stateful Processes

Since we updated our choreography model, we also need to update our process model to describe the implementations of choreographies. This is a straightforward extension of our previous calculus of simple processes, obtained by adding memories to processes. The new syntax is given in fig. 4.4.

A process term $\mathbf{p} \triangleright_v B$ now holds a value v , representing the memory state of the process. Send and receive actions are now extended to applying functions. A send action $\mathbf{p}!f$ sends the result of computing f in the local state of the process. Conversely, a receive action $\mathbf{p}?f$ computes f by using the value received from the other process and the local process memory, and then stores the result in the local process memory. An action f executes function f and updates the local memory of the process according to the result.

The semantics of stateful processes is also a straightforward extension, which uses the same evaluation function used for choreographies. The rules are given in fig. 4.5. The rules for the structural precongrence are the same (modulo the addition of values v in process terms, but they are influential), but we report them for the reader's convenience anyway in fig. 4.6.

4.1.3 EndPoint Projection

We have to update our definition of EPP to our new language model.

$$\begin{array}{c} \overline{(N_1 \mid N_2) \mid N_3 \equiv N_1 \mid (N_2 \mid N_3)} \text{ PA} \\ \overline{N_1 \mid N_2 \equiv N_2 \mid N_1} \text{ PC} \quad \overline{N \mid \mathbf{0} \preceq N} \text{ GCN} \quad \overline{\mathbf{p} \triangleright_v \mathbf{0} \preceq \mathbf{0}} \text{ GCP} \end{array}$$

Figure 4.6: Stateful processes, structural precongruence.

$$\begin{aligned} \llbracket \mathbf{p}.f \rightarrow \mathbf{q}.g; C \rrbracket_r &= \begin{cases} \mathbf{q}!f; \llbracket C \rrbracket_r & \text{if } r = \mathbf{p} \\ \mathbf{p}?g; \llbracket C \rrbracket_r & \text{if } r = \mathbf{q} \\ \llbracket C \rrbracket_r & \text{otherwise} \end{cases} \\ \llbracket \mathbf{p}.f; C \rrbracket_r &= \begin{cases} f; \llbracket C \rrbracket_r & \text{if } r = \mathbf{p} \\ \llbracket C \rrbracket_r & \text{otherwise} \end{cases} \\ \llbracket \mathbf{0} \rrbracket_{\mathbf{p}} &= \mathbf{0} \end{aligned}$$

Figure 4.7: Behaviour projection for stateful choreographies.

First, as usual, let us gain some intuition. Given any σ , the network implementation of the choreography given in example 10 should look like the following.

$$\begin{array}{c} \text{Buyer} \triangleright_{\sigma(\text{Buyer})} \text{Seller}!title; \text{Seller}?price; \mathbf{0} \\ | \\ \text{Seller} \triangleright_{\sigma(\text{Seller})} \text{Buyer}?x; \text{Buyer}!cat(x); \mathbf{0} \end{array}$$

We generalise this intuition in the following definition.

Definition 5 (EndPoint Projection (EPP)). *The EPP of a configuration $\langle C, \sigma \rangle$, denoted $\llbracket \langle C, \sigma \rangle \rrbracket$, is defined as:*

$$\llbracket \langle C, \sigma \rangle \rrbracket = \prod_{\mathbf{p} \in \text{procs}(C)} \mathbf{p} \triangleright_{\sigma(\mathbf{p})} \llbracket C \rrbracket_{\mathbf{p}}.$$

We also need to update the definition of behaviour projection— $\llbracket C \rrbracket_{\mathbf{p}}$ —since the language of stateful choreographies is different from that of simple choreographies. We display the new rules in fig. 4.7.

Appendix A

Solution to selected exercises

We give the solutions to some selected exercises. Some solutions are given in full detail, to serve as examples of exposition. All solutions should still provide enough information for the reader to figure out the remaining parts.

Solution of exercise 1. We prove only the direction from the system in fig. 1.1 to the system in fig. 1.2.

To prove this direction, we actually prove the stronger statement:

- $\text{conn}(A, B)$ provable in fig. 1.1 implies $\text{conn}(A, B)$ provable in fig. 1.2;
- $\text{path}(A, B)$ provable in fig. 1.1 implies $\text{path}(A, B, n)$ for some n provable in fig. 1.2.

The proof is by induction on the structure of the derivation of $\text{conn}(A, B)$ or $\text{path}(A, B)$. We proceed by cases on the last applied rule of the derivation.

Base cases (axioms) If the last applied rule is an axiom, then the proof is valid also in the other system.

Case Sym The derivation has this shape:

$$\frac{\frac{\mathcal{P}}{\text{conn}(B, A)}}{\text{conn}(A, B)} \text{SYM}.$$

By induction hypothesis, we know that there exists \mathcal{P}' in the other system such that:

$$\frac{\mathcal{P}'}{\text{conn}(B, A)}.$$

The thesis follows by applying rule SYM.

$$\frac{\frac{\mathcal{P}'}{\text{conn}(B, A)}}{\text{conn}(A, B)} \text{SYM}.$$

Case Dir The derivation has this shape:

$$\frac{\frac{\mathcal{P}}{\text{conn}(A, B)}}{\text{path}(A, B)} \text{DIR}.$$

By induction hypothesis, we know that there exists \mathcal{P}' in the other system such that:

$$\frac{\mathcal{P}'}{\text{conn}(A, B)}.$$

The thesis follows by applying rule DIRW.

$$\frac{\frac{\frac{\mathcal{P}'}{\text{conn}(A, B)}}{\text{path}(A, B, 1)}}{\text{path}(A, B, 1)} \text{DIRW}.$$

Case Trans The derivation has this shape:

$$\frac{\frac{\mathcal{P}}{\text{path}(A, B)} \quad \frac{\mathcal{Q}}{\text{path}(B, C)}}{\text{path}(A, C)} \text{TRANS}.$$

By induction hypothesis on \mathcal{P} and by induction hypothesis on \mathcal{Q} , we know that there exist n and m , \mathcal{P}' and \mathcal{Q}' in the other system such that:

$$\frac{\mathcal{P}'}{\text{path}(A, B, n)} \quad \frac{\mathcal{Q}'}{\text{path}(A, B, m)}.$$

The thesis follows by applying rule TRANSW.

$$\frac{\frac{\mathcal{P}'}{\text{path}(A, B, n)} \quad \frac{\mathcal{Q}'}{\text{path}(A, B, m)}}{\text{path}(A, B, n + m)} \text{TRANSW}.$$

■

List of Figures

1.1	An inference system for flights.	13
1.2	Weighted rules for flight paths.	17
1.3	A limited and weighted flight system.	18
1.4	An alternative way of constructing paths.	18
2.1	Simple choreographies, syntax.	24
2.2	Simple choreographies, semantics.	25
3.1	Simple processes, syntax.	28
3.2	Simple processes, structural relation.	32
3.3	Simple processes, semantics.	33
3.4	Behaviour projection for simple choreographies.	36
3.5	Simple concurrent choreographies, syntax.	38
3.6	Simple concurrent choreographies, semantics.	38
3.7	Simple concurrent choreographies, structural precongruence.	39
4.1	Stateful choreographies, syntax.	41
4.2	Stateful choreographies, semantics.	42
4.3	Stateful choreographies, structural precongruence.	43
4.4	Stateful processes, syntax.	45
4.5	Stateful processes, semantics.	45
4.6	Stateful processes, structural precongruence.	46
4.7	Behaviour projection for stateful choreographies.	46

LIST OF FIGURES

List of Notations

\mathcal{C} A context. 32

\mathcal{P} A derivation in an inference system. 16

Index

inference rule, 11
inference system, 11

Bibliography

Davide Ancona, Viviana Bono, Mario Bravetti, Joana Campos, Giuseppe Castagna, Pierre-Malo Deniélou, Simon J. Gay, Nils Gesbert, Elena Giachino, Raymond Hu, Einar Broch Johnsen, Francisco Martins, Viviana Mascardi, Fabrizio Montesi, Rumyana Neykova, Nicholas Ng, Luca Padovani, Vasco T. Vasconcelos, and Nobuko Yoshida. Behavioral types in programming languages. *Foundations and Trends in Programming Languages*, 3(2-3):95–230, 2016.

BPMN. Business Process Model and Notation. <http://www.omg.org/spec/BPMN/2.0/>, 2011.

Marco Carbone and Fabrizio Montesi. Deadlock-freedom-by-design: multi-party asynchronous global programming. In *POPL*, pages 263–274. ACM, 2013.

Marco Carbone, Kohei Honda, and Nobuko Yoshida. Structured communication-centred programming for web services. In Rocco De Nicola, editor, *ESOP*, volume 4421 of *LNCS*, pages 2–17. Springer, 2007. doi: 10.1007/978-3-540-71316-6_2.

Marco Carbone, Kohei Honda, and Nobuko Yoshida. Structured communication-centered programming for web services. *ACM Trans. Program. Lang. Syst.*, 34(2):8, 2012.

Luís Cruz-Filipe and Fabrizio Montesi. A core model for choreographic programming. In Olga Kouchnarenko and Ramtin Khosravi, editors, *FACS*, volume 10231 of *LNCS*. Springer, 2016. doi: 10.1007/978-3-319-57666-4_3.

Luís Cruz-Filipe and Fabrizio Montesi. Procedural choreographic programming. In *FORTE*, LNCS. Springer, 2017.

- Pierre-Malo Deniélou and Nobuko Yoshida. Multipart compatibility in communicating automata: Characterisation and synthesis of global session types. In *ICALP (2)*, pages 174–186, 2013.
- Nicola Dragoni, Saverio Giallorenzo, Alberto Lluch Lafuente, Manuel Mazara, Fabrizio Montesi, Ruslan Mustafin, and Larisa Safina. Microservices: yesterday, today, and tomorrow. In *Present and Ulterior Software Engineering*, pages 195–216. Springer, 2017.
- Xiang Fu, Tevfik Bultan, and Jianwen Su. Realizability of conversation protocols with message contents. *International Journal on Web Service Res.*, 2(4):68–93, 2005.
- Kohei Honda, Nobuko Yoshida, and Marco Carbone. Multipart Asynchronous Session Types. *J. ACM*, 63(1):9, 2016. doi: 10.1145/2827695. URL <http://doi.acm.org/10.1145/2827695>.
- Hans Hüttel, Ivan Lanese, Vasco T. Vasconcelos, Luís Caires, Marco Carbone, Pierre-Malo Deniélou, Dimitris Mostrous, Luca Padovani, António Ravara, Emilio Tuosto, Hugo Torres Vieira, and Gianluigi Zavattaro. Foundations of session types and behavioural contracts. *ACM Comput. Surv.*, 49(1):3:1–3:36, 2016.
- International Telecommunication Union. Recommendation Z.120: Message sequence chart, 1996.
- Ivan Lanese, Claudio Guidi, Fabrizio Montesi, and Gianluigi Zavattaro. Bridging the gap between interaction- and process-oriented choreographies. In *SEFM*, pages 323–332, 2008.
- Per Martin-Löf. On the meanings of the logical constants and the justifications of the logical laws. *Nordic journal of philosophical logic*, 1(1):11–60, 1996.
- Robin Milner. *A Calculus of Communicating Systems*, volume 92 of *LNCS*. Springer, Berlin, 1980.
- R.M. Needham and M.D. Schroeder. Using encryption for authentication in large networks of computers. *Commun. ACM*, 21(12):993–999, December 1978. ISSN 0001-0782. doi: 10.1145/359657.359659.
- OECD. Horizon scan of megatrends and technology trends in the context of future research policy, 2016. <http://ufm.dk/en/publications/2016/anoecd-horizon-scan-of-megatrends-and-technology-trends-in-the-context-of-future-research-policy>.

BIBLIOGRAPHY

- F. Pfenning. Lecture Notes on Deductive Inference, 2012.
<https://www.cs.cmu.edu/~fp/courses/15816-s12/lectures/01-inference.pdf>.
- Gordon D. Plotkin. A structural approach to operational semantics. *J. Log. Algebr. Program.*, 60-61:17–139, 2004.
- Z. Qiu, X. Zhao, C. Cai, and H. Yang. Towards the theoretical foundation of choreography. In *WWW*, pages 973–982. ACM, 2007.
- R. L. Rivest, A. Shamir, and L. Adleman. A method for obtaining digital signatures and public-key cryptosystems. *Commun. ACM*, 21(2):120–126, February 1978. ISSN 0001-0782. doi: 10.1145/359340.359342. URL <http://doi.acm.org/10.1145/359340.359342>.
- Davide Sangiorgi and David Walker. *The π -calculus: a Theory of Mobile Processes*. Cambridge University Press, 2001.